

Data-Cleaning-Report_Room-Temps

September 1, 2021

0.1 Imports

We'll import the *Pandas* and *mysql.connector* packages to import database data into a Pandas dataframe which we'll name **raw_temperature_df** (raw temperature dataframe). We'll also import a dictionary containing the login information to the MySQL server as **CREDS**.

Note that we'll leave **id** out of our SQL *SELECT* statement since Pandas provides automatic row indexing. These new indexes will be equal to the original **id** of the data minus 1.

```
[1]: import pandas as pd
import mysql.connector as connector
from database_credentials import MySQL_credentials as CREDS

connection = connector.connect(
    host = CREDS['host'],
    user = CREDS['user'],
    password = CREDS['password'],
    database = CREDS['database']
)

raw_temperature_df = pd.read_sql(f"SELECT inside_temperature,
    ↳outside_temperature, season, time, date FROM {CREDS['table']}",
    ↳con=connection)
```

0.2 Data Types

Next, we'll check the shape of the dataframe to confirm we imported all of the rows. We should have, at the time of writing, 347 rows. This is confirmed by accessing the *shape* attribute.

```
[2]: raw_temperature_df.shape
```

```
[2]: (347, 5)
```

We'll check that our columns, or attributes, are of the proper data types by accessing the *dtypes* attribute of the dataframe. In this case, the **date** column was incorrectly typed as an *object* (string) since Pandas doesn't support the *date* type that the **date** column was stored as in the database.

Note that Pandas also doesn't support the *time* type that the **time** column is stored as so it has been converted to a *timedelta*. This is fine for the purposes of our analysis.

```
[3]: raw_temperature_df.dtypes
```

```
[3]: inside_temperature      int64
      outside_temperature    int64
      season                 object
      time                   timedelta64[ns]
      date                   object
      dtype: object
```

To convert the **date** column to a *datetime* we'll use the `to_datetime` function that Pandas offers.

Another route to take would be re-querying the data, *CAST*ing the **date** column as *datetime* but that would be less computation efficient and less time efficient. This would also likely muddle the clarity of the data cleaning process.

Secondly, we'll create extra columns containing the year, month, and day of the month for each observation.

```
[4]: raw_temperature_df['date'] = pd.to_datetime(raw_temperature_df['date']) #
      ↪convert date column type from object to datetime

      raw_temperature_df['year'] = raw_temperature_df['date'].map(lambda x: x.year)
      raw_temperature_df['month'] = raw_temperature_df['date'].map(lambda x: x.month)
      raw_temperature_df['day'] = raw_temperature_df['date'].map(lambda x: x.day)
```

A final check of the data types in each column reveals exactly the desired outcome. The **year**, **month**, and **day** columns are as integers but this is fine for our purposes.

```
[5]: raw_temperature_df.dtypes
```

```
[5]: inside_temperature      int64
      outside_temperature    int64
      season                 object
      time                   timedelta64[ns]
      date                   datetime64[ns]
      year                   int64
      month                  int64
      day                    int64
      dtype: object
```

While data is only uploaded to the database if there are no null values in the observation, we should still take care to confirm this is the case for our dataframe. We use the *dropna* Pandas function to remove rows in the dataframe containing NaN values, the Pandas equivalent for null values, to be confident we won't raise any arithmetic exceptions during the analysis phase.

```
[6]: raw_temperature_df.dropna(axis='index', inplace=True)
```

0.3 Consistency of Data and Duplicates

We should make sure that the amount of data in each **time** group is consistent across groups to ensure that this doesn't introduce any potential biases. This isn't a necessary prerequisite if we wish to perform an [ANOVA](#) test, so long as the variance between the groups is similar, but it greatly improves the power of any statistical tests done. The power of a statistical test is only as strong as the group with the smallest sample size. For this reason we'll also check the variance of the groups by looking at the standard deviation for each.

Let's first start with the **time** groups before checking **day**, **month**, and **year**, which is also effectively season at this point.

```
[7]: raw_temperature_df.groupby(['time']).describe().loc[:, ['inside_temperature', 'outside_temperature']]
```

```
[7]:
```

		inside_temperature \					
		count	mean	std	min	25%	50%
time							
0 days	00:00:00	60.0	71.900000	2.790419	66.0	70.00	72.0
0 days	00:00:04	3.0	70.333333	2.081666	68.0	69.50	71.0
0 days	00:00:08	3.0	72.333333	0.577350	72.0	72.00	72.0
0 days	00:00:12	2.0	75.500000	0.707107	75.0	75.25	75.5
0 days	00:00:16	3.0	73.666667	1.527525	72.0	73.00	74.0
0 days	00:00:20	4.0	75.250000	2.500000	72.0	74.25	75.5
0 days	04:00:00	52.0	70.346154	1.866995	66.0	69.00	70.0
0 days	04:01:00	1.0	79.000000	NaN	79.0	79.00	79.0
0 days	08:00:00	53.0	70.830189	2.190327	66.0	69.00	71.0
0 days	08:01:00	1.0	70.000000	NaN	70.0	70.00	70.0
0 days	12:00:00	56.0	73.553571	2.682737	68.0	72.00	74.0
0 days	13:19:00	1.0	72.000000	NaN	72.0	72.00	72.0
0 days	16:00:00	53.0	73.981132	3.091493	67.0	72.00	74.0
0 days	20:00:00	54.0	73.333333	2.555054	67.0	72.00	74.0
0 days	20:01:00	1.0	81.000000	NaN	81.0	81.00	81.0


```

outside_temperature \
75%    max    count    mean    std    min
time
0 days 00:00:00 74.00 80.0    60.0 72.483333 4.575142 62.0
0 days 00:00:04 71.50 72.0    3.0 72.666667 1.527525 71.0
0 days 00:00:08 72.50 73.0    3.0 74.000000 2.645751 72.0
0 days 00:00:12 75.75 76.0    2.0 88.500000 2.121320 87.0
0 days 00:00:16 74.50 75.0    3.0 91.333333 3.785939 87.0
0 days 00:00:20 76.50 78.0    4.0 83.250000 3.685557 79.0
0 days 04:00:00 72.00 75.0    52.0 70.634615 4.401554 62.0
0 days 04:01:00 79.00 79.0    1.0 72.000000      NaN 72.0
0 days 08:00:00 72.00 76.0    53.0 72.037736 5.045799 61.0
0 days 08:01:00 70.00 70.0    1.0 68.000000      NaN 68.0
0 days 12:00:00 75.00 79.0    56.0 83.267857 6.934738 67.0

```

0 days 13:19:00	72.00	72.0		1.0	77.000000	NaN	77.0
0 days 16:00:00	76.00	81.0		53.0	82.207547	6.805989	67.0
0 days 20:00:00	75.00	79.0		54.0	77.648148	5.508681	65.0
0 days 20:01:00	81.00	81.0		1.0	74.000000	NaN	74.0

	25%	50%	75%	max
time				
0 days 00:00:00	69.75	73.0	75.25	82.0
0 days 00:00:04	72.00	73.0	73.50	74.0
0 days 00:00:08	72.50	73.0	75.00	77.0
0 days 00:00:12	87.75	88.5	89.25	90.0
0 days 00:00:16	90.00	93.0	93.50	94.0
0 days 00:00:20	82.00	83.0	84.25	88.0
0 days 04:00:00	68.00	72.0	73.25	80.0
0 days 04:01:00	72.00	72.0	72.00	72.0
0 days 08:00:00	69.00	72.0	75.00	85.0
0 days 08:01:00	68.00	68.0	68.00	68.0
0 days 12:00:00	78.00	83.0	89.25	95.0
0 days 13:19:00	77.00	77.0	77.00	77.0
0 days 16:00:00	78.00	81.0	87.00	97.0
0 days 20:00:00	74.00	76.5	81.00	91.0
0 days 20:01:00	74.00	74.0	74.00	74.0

The first thing to notice is that, there are some observations that were recorded at odd times. This could be caused due to a power outage delaying the running of the data collection script, an error in uploading to the database, or simple network/system lag. We can allow for a margin of error of a minute, since temperature changes are typically negligible on the timescales of seconds. Upon further examination, the discrepancy between the count of observations at midnight (00:00:00) and the other times seems a bit conspicuous, especially if we were to assign the times between 00:00:04 and 00:00:20 to also be midnight observations.

A more meticulous manual look at the data shows an interesting pattern between indices 64 and 80 and it becomes clear what occurred.

```
[8]: raw_temperature_df.loc[62:83, ['time', 'date']]
```

```
[8]:
```

	time	date
62	0 days 08:00:00	2021-07-06
63	0 days 12:00:00	2021-07-06
64	0 days 00:00:16	2021-07-06
65	0 days 00:00:20	2021-07-06
66	0 days 00:00:08	2021-07-07
67	0 days 00:00:12	2021-07-07
68	0 days 00:00:16	2021-07-07
69	0 days 00:00:20	2021-07-07
70	0 days 00:00:04	2021-07-08
71	0 days 00:00:08	2021-07-08

```

72 0 days 00:00:12 2021-07-08
73 0 days 00:00:20 2021-07-08
74 0 days 00:00:00 2021-07-09
75 0 days 00:00:04 2021-07-09
76 0 days 00:00:16 2021-07-09
77 0 days 00:00:20 2021-07-09
78 0 days 00:00:00 2021-07-10
79 0 days 00:00:04 2021-07-10
80 0 days 00:00:08 2021-07-10
81 0 days 12:00:00 2021-07-10
82 0 days 16:00:00 2021-07-10
83 0 days 20:00:00 2021-07-10

```

From the latter half of 2021-07-06 through the first half of 2021-07-10, all times were formatted incorrectly and thus we can adjust them using the *replace* Pandas function. We'll then use the same function to adjust the 04:01:00, 08:01:00, and 20:01:00 timed observations. We'll also create a new dataframe **fixed_times_df** to move forward.

Finally, we'll use the *drop* function to eliminate the observation recorded at 13:19:00 and use *drop_duplicates* based on **date** and **time** to eliminate any conflicting observations and repeated entries.

```

[9]: fixed_times_df = raw_temperature_df.replace(
      to_replace=pd.to_timedelta(['00:00:04', '00:00:08', '00:00:12', '00:00:16',
      ↪ '00:00:20']),
      value=pd.to_timedelta(['04:00:00', '08:00:00', '12:00:00', '16:00:00', '20:
      ↪ 00:00']),
      )

fixed_times_df.replace(
      to_replace=pd.to_timedelta(['04:01:00', '08:01:00', '20:01:00']),
      value=pd.to_timedelta(['04:00:00', '08:00:00', '20:00:00']),
      inplace=True
      )

fixed_times_df.drop(fixed_times_df[fixed_times_df['time'] == pd.
      ↪ to_timedelta('13:19:00')].index, inplace=True)

fixed_times_df.drop_duplicates(['date', 'time'], keep='first', inplace=True)

```

Checking again on the description of the dataframe, grouped by time, we now have more consistent group sizes and standard deviations (a measure of variance) across groups.

```

[10]: fixed_times_df.groupby(['time']).describe().loc[:, ['inside_temperature',
      ↪ 'outside_temperature']]

```

```

[10]:
      inside_temperature
count      mean      std  min  25%  50%  \

```

time								
0 days 00:00:00	60.0	71.900000	2.790419	66.0	70.0	72.0		
0 days 04:00:00	56.0	70.500000	2.174229	66.0	69.0	70.0		
0 days 08:00:00	57.0	70.894737	2.143797	66.0	69.0	71.0		
0 days 12:00:00	58.0	73.620690	2.661141	68.0	72.0	74.0		
0 days 16:00:00	56.0	73.964286	3.020923	67.0	72.0	74.0		
0 days 20:00:00	59.0	73.593220	2.736211	67.0	72.0	74.0		

	outside_temperature					
	75%	max	count	mean	std	min
time						
0 days 00:00:00	74.0	80.0	60.0	72.483333	4.575142	62.0
0 days 04:00:00	72.0	79.0	56.0	70.767857	4.276749	62.0
0 days 08:00:00	72.0	76.0	57.0	72.070175	4.938405	61.0
0 days 12:00:00	75.0	79.0	58.0	83.448276	6.885460	67.0
0 days 16:00:00	76.0	81.0	56.0	82.696429	6.972464	67.0
0 days 20:00:00	75.0	81.0	59.0	77.966102	5.542830	65.0

	25%	50%	75%	max
time				
0 days 00:00:00	69.75	73.0	75.25	82.0
0 days 04:00:00	68.00	72.0	73.25	80.0
0 days 08:00:00	69.00	72.0	75.00	85.0
0 days 12:00:00	78.25	83.0	89.75	95.0
0 days 16:00:00	78.75	82.0	87.00	97.0
0 days 20:00:00	74.00	77.0	81.50	91.0

We'll check the same, now grouping by **day**, **month**, and **year** (which will have the same output as **season** since data collection has only occurred during summer). All of these have consistent standard deviations and ranges. There is some inconsistency between **day** groups and **month** groups but, because the variance in the temperature readings is similar, we can proceed. It should be noted, however, that any statistical tests run will only have the power based on the smallest group in the group pool.

```
[11]: fixed_times_df.groupby(['day']).describe().loc[:, ['inside_temperature',
→ 'outside_temperature']]
```

	inside_temperature							
	count	mean	std	min	25%	50%	75%	max
day								
1	12.0	71.583333	2.539088	69.0	69.00	72.0	73.00	76.0
2	12.0	71.250000	2.632835	68.0	69.00	71.0	73.00	76.0
3	10.0	69.000000	1.490712	67.0	68.00	68.5	70.00	72.0
4	11.0	69.909091	2.300198	66.0	68.50	70.0	71.00	74.0
5	9.0	70.666667	2.121320	67.0	70.00	70.0	72.00	74.0
6	11.0	73.181818	2.993933	70.0	71.50	72.0	74.00	79.0
7	9.0	73.444444	2.006932	71.0	72.00	72.0	75.00	76.0

8	10.0	72.800000	1.686548	71.0	72.00	72.0	73.75	76.0
9	10.0	72.000000	3.162278	68.0	70.00	71.0	74.00	78.0
10	11.0	72.454545	2.910795	68.0	70.00	73.0	75.00	76.0
11	7.0	70.857143	2.794553	68.0	69.00	70.0	72.00	76.0
12	10.0	74.100000	2.726414	71.0	72.00	73.0	76.50	78.0
13	11.0	74.090909	2.700168	70.0	72.50	74.0	75.00	79.0
14	11.0	75.454545	2.252272	72.0	73.50	77.0	77.00	78.0
15	12.0	72.333333	3.200379	67.0	71.00	72.5	73.50	79.0
16	11.0	73.454545	1.967925	70.0	72.50	74.0	75.00	75.0
17	9.0	74.222222	3.032234	71.0	73.00	74.0	75.00	81.0
18	10.0	73.200000	1.988858	69.0	72.25	74.0	74.75	75.0
19	10.0	72.100000	2.024846	70.0	71.00	71.5	72.00	77.0
20	10.0	71.600000	2.366432	69.0	70.25	71.0	72.75	77.0
21	9.0	70.000000	1.224745	68.0	70.00	70.0	71.00	71.0
22	10.0	72.600000	2.170509	70.0	71.00	71.5	74.75	76.0
23	10.0	72.400000	1.712698	69.0	72.00	72.0	73.75	75.0
24	11.0	71.363636	2.203303	69.0	70.00	70.0	72.50	75.0
25	17.0	71.294118	3.235829	66.0	70.00	71.0	74.00	78.0
26	15.0	70.933333	3.104528	66.0	69.00	70.0	73.50	76.0
27	14.0	72.214286	3.166618	66.0	71.00	72.0	74.00	78.0
28	14.0	73.000000	2.855494	68.0	71.00	73.0	74.00	78.0
29	15.0	74.466667	3.833437	68.0	72.00	75.0	77.00	81.0
30	16.0	74.687500	2.868652	70.0	72.75	75.0	75.50	80.0
31	9.0	72.888889	2.420973	70.0	71.00	72.0	74.00	78.0

	outside_temperature							
	count	mean	std	min	25%	50%	75%	max
day								
1	12.0	73.750000	6.877169	62.0	70.25	75.0	77.00	87.0
2	12.0	70.916667	5.107184	63.0	69.25	71.0	73.25	79.0
3	10.0	67.800000	4.732864	63.0	64.25	66.0	70.25	77.0
4	11.0	70.000000	6.115554	63.0	64.50	70.0	74.50	80.0
5	9.0	73.333333	6.670832	64.0	68.00	73.0	79.00	83.0
6	11.0	78.000000	9.695360	66.0	70.00	76.0	84.50	93.0
7	9.0	79.333333	9.526279	68.0	72.00	77.0	88.00	94.0
8	10.0	72.800000	7.524774	65.0	67.00	70.5	76.25	87.0
9	10.0	75.400000	6.535374	67.0	70.75	74.0	79.50	87.0
10	11.0	75.727273	5.159281	71.0	71.50	74.0	79.50	85.0
11	7.0	77.857143	6.914443	71.0	72.50	76.0	81.50	90.0
12	10.0	81.500000	8.168367	74.0	75.25	77.0	87.00	95.0
13	11.0	79.727273	7.988628	72.0	73.50	78.0	84.50	94.0
14	11.0	81.363636	6.407382	72.0	77.00	82.0	85.50	91.0
15	12.0	78.500000	5.916080	72.0	74.00	76.0	83.25	88.0
16	11.0	80.181818	7.820718	70.0	74.50	80.0	84.50	93.0
17	9.0	80.222222	4.918785	74.0	77.00	80.0	81.00	91.0
18	10.0	77.900000	5.173651	71.0	74.00	77.5	82.50	85.0
19	10.0	78.100000	6.190495	70.0	72.25	79.0	81.50	89.0

20	10.0	76.300000	5.396501	69.0	72.50	75.0	79.25	87.0
21	9.0	75.888889	4.859127	70.0	73.00	74.0	78.00	86.0
22	10.0	73.000000	5.077182	64.0	71.25	73.0	75.00	81.0
23	10.0	73.100000	6.118279	64.0	68.25	73.0	77.50	83.0
24	11.0	75.454545	7.257598	64.0	70.50	74.0	80.50	89.0
25	17.0	75.647059	8.659830	62.0	71.00	74.0	80.00	91.0
26	15.0	78.266667	6.419464	69.0	73.50	77.0	82.00	92.0
27	14.0	80.142857	7.998626	68.0	74.00	77.5	87.00	94.0
28	14.0	78.071429	7.710775	69.0	73.00	75.5	83.50	93.0
29	15.0	76.800000	8.945869	65.0	70.50	75.0	78.00	95.0
30	16.0	80.312500	8.584628	70.0	72.00	79.5	85.25	97.0
31	9.0	74.777778	7.980880	61.0	72.00	75.0	78.00	86.0

```
[12]: fixed_times_df.groupby(['month']).describe().loc[:, ['inside_temperature',
↪ 'outside_temperature']]
```

```
[12]:
```

	inside_temperature							
	count	mean	std	min	25%	50%	75%	max
month								
6	34.0	70.529412	3.057383	66.0	68.0	70.5	73.0	79.0
7	157.0	72.254777	2.700675	66.0	70.0	72.0	74.0	79.0
8	155.0	73.000000	2.956393	67.0	71.0	73.0	75.0	81.0

	outside_temperature							
	count	mean	std	min	25%	50%	75%	max
month								
6	34.0	79.823529	9.440309	62.0	74.0	78.0	87.75	97.0
7	157.0	76.286624	7.227131	61.0	71.0	75.0	81.00	94.0
8	155.0	76.129032	7.330518	62.0	71.0	74.0	80.00	95.0

```
[13]: fixed_times_df.groupby(['year']).describe().loc[:, ['inside_temperature',
↪ 'outside_temperature']]
```

```
[13]:
```

	inside_temperature							
	count	mean	std	min	25%	50%	75%	max
year								
2021	346.0	72.419075	2.934174	66.0	70.0	72.0	74.0	81.0

	outside_temperature							
	count	mean	std	min	25%	50%	75%	max
year								
2021	346.0	76.563584	7.569409	61.0	71.0	75.0	81.0	97.0

0.4 Conclusion

The data is now sufficiently cleaned and may be used for analysis.


```
[14]: clean_df = fixed_times_df
```

```
[15]: clean_df.to_csv('../cleaned_data.csv')
```